

BINARY SEARCH TREES (CONTD)

Problem Solving with Computers-II

C++

```
#include <iostream>
using namespace std;

int main(){
    cout<<"Hola Facebook\n";
    return 0;
}
```

How is PA02 going?


- A. Done!
- B. On track to finish
- C. On track to finish but my code is a mess
- D. Stuck and struggling
- E. Haven't started

Midterm – Monday 2/26

- We will have a review session (by our TA Jack) on Friday 1-5p in Phelps 3526
- Which of the following review session times work best for you?
 - A. 1 - 3p
 - B. 3p - 5p
 - C. Neither

Please post the topics that you need help with on Piazza.

Binary Search Tree – Review



- Rooted binary tree
- Each node:
 - stores a key
 - has a pointer to left child, right child and parent (optional)
- Satisfies the **Search Tree Property**


Height of the tree



Many different BSTs are possible for the same set of keys
Examples for keys: 12, 32, 41, 42, 45


- *What is the height in each case?*
- *How do we define the height of the tree?*

Balanced trees



- What is a balanced tree?


Relating H (height) and N (#nodes)
find is $O(H)$, we want to find a $f(N) = H$



How many nodes are on level L in a completely filled binary search tree?

- A. 2
- B. L
- C. $2 * L$
- D. 2^L

Relating H (height) and N (#nodes)
find is $O(H)$, we want to find a $f(N) = H$




Finally, what is the height (exactly) of the tree in terms of N ?

Binary Search Trees

- WHAT are the operations supported?
- HOW do we implement them?
- WHY do we get certain run times?

Search



- Search for 41, then 40


How many pointer traversals are needed to search for 40?

- A. One
- B. Two
- C. Three
- D. Four


Searching for a Key

Alg: TREE-SEARCH(x, k)

1. **if** $x = \text{NIL}$ or $k = \text{key}[x]$
2. **then return** x
3. **if** $k < \text{key}[x]$
4. **then return** TREE-SEARCH(left $[x], k$)
5. **else return** TREE-SEARCH(right $[x], k$)



Insert




- Insert 40
- Search for the key
- Insert at the spot you expected to find it

What is the worst case running time of search or insert?

- A. $O(1)$
- B. $O(\log N)$
- C. $O(\text{height})$
- D. $O(N)$

Min: Search for negative infinity



Can you come up with a recursive implementation for min?

A. Yes

B. No, it has to be iterative


Finding the Minimum in a Binary Search Tree

Goal: find the minimum value in a BST

Following left child pointers from the root, until a NIL is encountered


Alg: TREE-MINIMUM(x)

1. **while** left [x] \neq NIL
2. **do** x \leftarrow left [x]
3. **return** x



Minimum = 2


Max: Search for positive infinity



Traversing a Binary Search Tree

- **Inorder** tree walk:
 - Root is printed between the values of its left and right subtrees:
left, root, right → keys are printed in **sorted order**
- **Preorder** tree walk: root printed first: **root, left, right**
- **Postorder** tree walk: root printed last: **left, right, root**

In order traversal: print elements in sorted order




What is the running time?

Traversing a Binary Search Tree

Alg: **INORDER-TREE-WALK(x)**

1. **if** $x \neq \text{NIL}$
2. **then** INORDER-TREE-WALK (left [x])
3. print key [x]
4. INORDER-TREE-WALK (right [x])


E.g.:




Output: 2 3 4 **5** 7 9

- Running time:
 - $\Theta(N)$, where N is the size of the tree rooted at x

Pre-order traversal: What is it good for?



Predecessor: Next smallest element




- What is the predecessor of 32?
- What is the predecessor of 45?

Worst case analysis

Are binary search trees *really* faster than linked lists for finding elements?

- A. Yes
- B. No




Summary of operations

Average case analysis of a “successful” find

Given a BST having N nodes x_1, \dots, x_N , such that $\text{key}(x_i) = k_i$

How many compares to locate a key in the BST?

1. Worst case:
2. Best case:
3. Average case:



Here is the result! Proof is a bit involved but if you are interested in the proof, come to office hours

 $D_{avg}(N)$

Average #comparisons to find a single item in any BST with N nodes

$$D_{avg}(N) \approx 1.386 \log_2 N$$

Conclusion: The average time to find an element in a BST with no restrictions on shape is $\Theta(\log N)$.